

The Impact of Outcome Preferences in a Collection of Non-Zero-Sum Grid Games

Joseph L. Austerweil[†] and Stephen Brawner^{*} and Amy Greenwald^{*} and Elizabeth Hilliard^{*} and Mark Ho^{†*} and Michael L. Littman^{*} and James MacGlashan^{*} and Carl Trimbach^{*}

Brown University

^{*}Department of Computer Science

[†]Department of Cognitive, Linguistic, and Psychological Sciences
Providence, RI 02912

Abstract

We examined the behavior of reinforcement-learning algorithms in a set of two-player stochastic games played on a grid. These games were selected because they include both cooperative and competitive elements, highlighting the importance of adaptive collaboration between the players. We found that pairs of learners were surprisingly good at discovering stable mutually beneficial behavior when such behaviors existed. However, the performance of learners was significantly impacted by the subjective reward functions of the players. We found similar patterns of results in games involving human-human and human-agent pairs.

The field of reinforcement learning (Sutton and Barto 1998) is concerned with agents that improve their behavior in sequential environments through interaction. One of the best known and most versatile reinforcement-learning (RL) algorithms is Q-learning (Watkins and Dayan 1992), which is known to converge to optimal decisions in environments that can be characterized as Markov decision processes. Q-learning is best suited for single-agent environments; nevertheless, it has been applied in multi-agent environments (Sandholm and Crites 1995; Gomes and Kowalczyk 2009; Wunder, Littman, and Babes 2010), including non-zero-sum stochastic games, with varying degrees of success.

Nash-Q (Hu and Wellman 2003) is an attempt to adapt Q-learning to the general-sum setting, but its update rule is inefficient and it lacks meaningful convergence guarantees (Bowling 2000; Littman 2001). Correlated-Q (Greenwald and Hall 2003) is an improvement over Nash-Q in that, in exchange for access to a correlating device, its update rule is computationally efficient. However, there exist environments in which correlated-Q does not converge to optimal decisions (Zinkevich, Greenwald, and Littman 2005). Minimax-Q (Littman 1994a) converges to provably optimal decisions, but only in zero-sum Markov games. Likewise, Friend-Q and Foe-Q (Littman 2001) provably converge, but only to optimal decisions in purely cooperative and purely competitive games, respectively.

One significant shortcoming of the aforementioned multi-agent learning algorithms is that they define their updates in

a way that makes assumptions about their opponents without actually factoring in the opponents' observed behavior. In a sense, they are too stubborn. In contrast, single-agent learning algorithms like Q-learning are too flexible—they simply adapt to their opponents without consideration of how their behavior will impact the opponent. What is lacking in these existing algorithms is the ability to *negotiate* a mutually beneficial outcome (Gal et al. 2004).

Algorithms have been designed that seek a best response against a fixed player and a mutually beneficial response against like players (Conitzer and Sandholm 2007; Bowling and Veloso 2002). Others attempt to “lead” a learning opponent to beneficial behavior (Littman and Stone 2001). In this work, we return to the investigation of the behavior of single-agent Q-learning in multi-agent environments.

Our main contribution is to introduce a form of negotiation by incorporating into an agent's world view other-regarding preferences. Our approach goes beyond earlier attempts to nudge agents toward more cooperative behavior (Babes, Munoz de Cote, and Littman 2008) by providing a general framework that separates objective and subjective rewards (Singh et al. 2010) with the goal of increasing objective reward. We investigate the behavior of this approach in machine-machine and machine-human interactions.

Experimental Testbed

For our work, we devised and adapted several two-agent grids that require differing levels of coordination, and also allow agents to defend against uncooperative partners. On each turn of a round in our grid games, agents choose one of five actions (north, south, east, west, wait), which are then executed simultaneously. Agent transition dynamics are deterministic, and there is no tie-breaking when two agents collide. Instead, both agents remain on their current location if their chosen actions would result in a collision with one another. A round ends when either (or both) players move into their goal or a maximum number of turns has been reached.

The three-by-five grid in Figure 1 (Hallway) allows both agents to coordinate through a strategy where one agent moves along the top row and the other along the bottom row without interfering with one another. However, an agent that notices that the other is stepping aside might choose to play by “defecting” and proceeding straight to the goal. There are




Figure 1: Hallway, a three-by-five grid that requires an agreed-upon coordination strategy to efficiently play the game. Several coordination strategies also allow the agent to defend against an uncooperative partner without losing the round.




Figure 2: Intersection, a three-by-five grid that requires Blue to defend Orange's goal to encourage Orange to cooperate and move to the end of the uppermost hallway.

strategies to defend against such a non-cooperative agent, however. For example, if Orange moves south initially to (1, 3) and Blue moves west to (4, 2), Orange might choose to return and remain on its goal until Blue retreats to (4, 3) or (4, 1), at which point the players are equidistant from their goals and can continue safely from there. A more explicit defensive strategy for Orange would have it proceed east to (2, 3), and then choose to move north to (2, 2). If Blue is in (3, 2) and attempts to move west into (2, 2), the agents would collide until one agent chooses a different action. It is in Orange's best interest to continue choosing the north action to try to take (2, 2) and wait until Blue surrenders and chooses to go north into (3, 3) to avoid collision. At this point, the agents are equidistant from their goals and can continue safely from there.

The grid in Figure 2 (Intersection) requires Blue to defend against the possibility of an uncooperative orange agent by squatting on the orange goal. Orange should then move to (3, 1) where both agents are equidistant from their goals.

Figure 3 (Door) is a grid that requires coordination to navigate through the narrow center space at (3, 2).

In the grid in Figure 4 (Long hall), Blue begins one step closer to its goal than does Orange. However, Orange can squat on the blue goal until Blue chooses to cooperate by stepping one square back. If Orange can predict when Blue steps back, Orange can pivot left and step one closer to the goal while Blue attempts to step further away.

Our last grid, shown in Figure 5 (No compromise), requires both agents to trust and cooperate with each other to




Figure 3: Door, a three-by-five grid that requires the partners to agree on an order to navigate through the center square.



Figure 4: Long hall, a one-by-seven grid that allows Orange to squat on Blue's goal should they choose to not cooperate.

arrive at the goals at the same time. For example, Orange may sit on Blue's goal to let it pass into square (1, 2). Then, Blue must wait two turns before both agents are equidistant from the goals. If Blue defects and moves north into (1, 1) while Orange moves south from (2, 1) into (2, 2), Orange has the opportunity to step north to block Blue from reaching the goal. However, if Blue moves north into (1, 1) when Orange steps east into (3, 2), Blue will arrive at its goal sooner. Therefore, a trust spanning multiple rounds is required for the agents to effectively cooperate.

Machine-Machine Experiments

We carried out a set of simulation experiments with Q-learning in the grids described in the previous section.

Q-Learning in Self-Play

For each grid, we conducted 50 independent runs in which two Q-learners faced off. Runs lasted for 30,000 rounds each. To ensure that the state space was adequately explored, any state that is feasibly reachable from the default initial state had some probability of being selected as the starting position for a round. The agents' value functions were optimistically initialized with a value of 40 for all states and they used Boltzmann exploration with a temperature of 0.5. The agents were not charged any step costs and received objective rewards of 50 upon reaching their respective goals.




Figure 5: No compromise, a two-by-five grid that requires one agent to sit on the other's goal and then move to its goal when the other agent has passed.




Figure 6: Fraction of runs in which Q-learning agents ended with a cooperative strategy in which both agents reached the goal.

Once any agent reached its goal (or 100 moves were taken), the round was terminated. Agents used a discount factor 0.9 and learning rate of 0.01.

We denote the outcome of a round as a pair of letters, where **G** means the agent reached the goal and **N** means the agent did not reach the goal. The first letter in the pair represents the agent’s own outcome and the second represents the opponent’s outcome. For example, **GN** is used to denote that the agent reaches its goal while the opponent does not.

At the end of the 30,000 rounds, we checked the policies the learners had constructed. Figure 6 shows how often the learners found mutually beneficial cooperative policies that lead to **GG** outcomes. Note that, in spite of the learning algorithm not explicitly seeking outcomes with high social welfare, it very reliably identified such policies.

As is visible in Figure 6, only No compromise posed a challenge to the learners. Q-learning here tends to thrash about, finding a pair of policies that work well together only to eventually discover that one of the agents has an opportunity to defect. The defection destabilizes the policies and restarts the search for policies that work well together. Sometimes, Q-learning finds policies that are cooperative and stable, but these policies are brittle and would not be successful against agents with other strategies.

In the next section, we propose a classification scheme for policies and identify an important property that helps explain why some grids supported consistent cooperation and others did not. The classification also provides insight into what “better” approaches to learning might look like.

Cooperatively Defensive Policies

In the description of Hallway, we noted the existence of a pair of policies that agents could execute that would allow both agents to reach their goals while simultaneously preventing defecting—an agent reaching the goal alone. Figure 7 illustrates one such policy pair for this grid.

We call strategies that allow each agent to choose a mutually agreeable sequence of actions while also defending against an agent that is uncooperative as *cooperatively defensive* (CD) strategies. Formally, policy π is C (cooperative) if there exists an opponent policy π' such that **GG** is a possible outcome of π vs. π' , and a policy π is D (defensive) if there does not exist an opponent policy π' such that **NG** is




Figure 7: An illustration of a CD policy for Hallway—Blue allows Orange to score but only Orange if takes a detour that ensures that it doesn’t reach the goal first.

a possible outcome of π vs. π' . A CD policy is one that it is both C and D.

There are several things to note about CD policies. First, an agent can enter its goal first (**GN**) and still have played a CD policy. For the policy to be CD, an opponent policy with outcome **GG** has to exist. However, the opponent may not play a policy such that it reaches a goal. Second, both agents could not end up in their goal (**NN**). This lose-lose outcome could occur if the opponent’s policy is purely defensive (for example, stand on the agent’s goal) and it could even happen if the two agents play “mismatched” CD policies. In Hallway, a mismatch could happen if both agents go north and then end up running into each other in the middle square of the top row.

Briefly, here is an analysis of CD strategies for each of the five test grids.

- **Hallway:** As noted, there are several CD strategies in Hallway. The key feature of a CD strategy in this grid is that the player does not move more than one step from its goal until its opponent has taken an “extra” step (either waiting or stepping north or south). For the strategy to allow for the opponent to also play a CD strategy, the first move must be to wait or take a step north or south.
- **Intersection:** Purely cooperative agents could adopt a strategy in which Orange waits a single step before proceeding into its goal. This strategy is not CD, however, since Blue does not have the opportunity to observe if Orange will cooperate (wait) or defect (go north) and therefore cannot defend against Orange if it heads straight toward its goal. In contrast, the strategy described earlier for this grid in which Blue sits on Orange’s goal requires more actions. However, it is CD as both players are able to reach a position in which they are three turns from their goal without ever letting the other agent get closer than it is.

- **Door:** The simplest CD strategy for this grid is asymmetric and requires one agent to cede to the other agent the center square. For example, if Orange chooses to cede that space, it should step west into (2, 3) while Blue steps south into (3, 2). Then, Orange needs to step east back into (3, 3) to prevent Blue from marching straight for the goal. Only when Blue agrees to step aside to (2, 2) will they both be equidistant from their respective goals and in position to cooperate.
- **Long hall:** The strategy that minimizes the risk to either agent requires that Blue wait one turn initially, while Orange moves toward its goal. It is CD.
- **No compromise:** No pair of CD strategies exists for this grid. In particular, it is like Door in that only one player can go through the middle square at a time. Unlike Door, however, there's no way for both agents to stay equally far from their goals. As a result, there's always an incentive for one agent to defect and there's nothing the agent can do to stop it. Note however, that if Orange sits on the blue goal while Blue walks to (1, 1) and then Blue waits until Orange is at (1, 3) to move toward its goal, Blue's policy is technically CD while Orange's is C. Orange does not have a CD policy in response to Blue's CD policy.

Whereas Q-learning can find, and converge to, a stable, mutually beneficial set of strategies in four of the grids, doing so is not possible in No compromise. Q-learning instead tends to thrash about, finding a pair of policies that work well together only to eventually discover that one of the agents has an opportunity to defect, destabilizing the policies and prolonging the search indefinitely.

Whereas Q-learning can find, and converge to, stable, mutually beneficial strategies in four of the grids, the policies found are not always CD. In the next section, we examine a modification of Q-learning that performs more reliably across all five grids.

Subjective Rewards

In this section, we examine the impact of endowing agents with *other-regarding preferences*. That is, their rewards will no longer necessarily depend solely upon the world as it pertains to themselves alone. They also consider the effects of their actions on the outcomes of the other agents operating in the world. To make this idea more precise, we separate *objective* and *subjective* reward. Objective reward is simply the reward signal provided to the agent from the environment and by which behavior is judged. Standard reinforcement-learning agents, such as a Q-learning agent, seek to optimize their own objective reward signal—we call this preference the “selfish” preference because these agents are only concerned with their own outcomes. In some environments, however, it useful to distinguish this objective reward from subjective reward—the quantity the agent *believes* it should be optimizing. Previous work has shown that optimizing subjective reward can sometimes lead agents to be more effective in their acquisition of objective reward (Singh, Lewis, and Barto 2009).

Considering the 4 different outcomes in these games—GG, GN, NG, NN—there are 75 different possi-




Figure 8: Average score in self play after 30,000 rounds.

ble preference orderings when allowing for ties (see <https://oeis.org/A000670>). The selfish ordering that ignores the opponent's outcome is one of these. We write the ordering $GG \sim GN \succ NG \sim NN$. That is, regardless of the opponent's outcome, a selfish agent only prefers that it gets to its goal. Of the 75 distinct preference orderings, 9 of them are consistent with the selfish ordering, strictly preferring Gx to Nx for all x .

Of particular interest is the *fair* preference, defined to be the objective reward of the agent minus 25% of the difference between its own and the opponent's objective rewards:

$$r_s = r_a - 0.25 |r_a - r_o|,$$

where r_s is the agent's subjective reward, r_a is the agent's objective reward, and r_o is the opponent's objective reward.¹ By incorporating this fairness term into the agent's subjective reward, it strictly prefers the following order of outcomes:

$$GG \succ GN \succ NN \succ NG. \quad (1)$$

That is, the agent prefers making it to its goal as opposed to not, but it additionally prefers that the opponent only get to its goal if the agent itself does as well. To say it another way, the fair agent wants its opponent to win with it or lose with it.

Figure 8 shows the result of the selfish and fair agents playing against others of the same type in each of our test grids. In all five grids, the fair agents are able to obtain higher total (objective) reward.

Of the 9 orders, only 3 (all variations of the fair preference orderings) achieve consistent cooperation in self play across all 5 grids. In addition, the average score (across both players) for games involving the fair preference ordering is higher than any other preference ordering.

In Figure 9, we show results on the classification of policies found by the agents when learning against opponents of different preference types. These results show that the fair agent (Q-learning with preferences in Equation 1) finds a CD policy as frequently as or more frequently than normal selfish Q-learning when trained against another Q-learning agent or when trained against a fair agent. The exception is Blue in Intersection. Across all five games and all agent

¹Other percentages would achieve the same result.




Figure 9: The percentage of learning attempts where the agent’s final policy was CD by game and learning opponent.

positions, the selfish agent finds a CD policy in 50.9% attempts when learning against selfish and in 12.8% of attempts against the fair agent. The fair agent’s performance in these cases is 88% and 25.5%, respectively. These results suggest that Q-learners with fair preferences may find CD policies more often, especially when they learn against a selfish player. Whether this is general for a broader set of grids is an important question for future work.

Given the option of who to play with, agents do best when playing against the fair preference ordering. We also analyzed which preference ordering is the preferred one to adopt. It turns out that the preference ordering $\text{GN} \succ \text{GG} \succ \text{NG} \sim \text{NN}$ is an evolutionary stable strategy—it outperforms competitors in a population and can maintain its dominance against invaders. It adopts a much more aggressive stance than the fair preference ordering in that it prefers winning alone to sharing the top spot. Thus, it has a tendency to find strategies that drives its opponents’ scores down.

Human–Human Experiments

Ultimately, a major goal for developing machine agents that act intelligently in multi-agent scenarios is to apply them to real world problems. Humans are already agents that machine agents interact with in some current multi-agent environments (such as the stock market and online advertising auctions). Successfully expanding the scope of applications where multi-agent learning can be applied in the real world necessitates studying how these agents interact with human agents. A machine agent that interacts optimally against other machine agents, but not against human agents, is not likely to be effective in environments that include human agents. Further, one major goal of developing machine agents is for them to solve tasks in collaboration with human agents. Given the controversial nature of rationality assumptions for human agents (Kahneman, Slovic, and Tversky 1982), a machine agent that plans its collaboration by assuming the human agent will act rationally (optimally) is unlikely to be successful in collaborating with the human agent. Thus, in this section, we investigate how human agents interact with each other in Hallway and in the next section, how human agents interact with fair and selfish learning agents in Hallway.

A total of 40 human participants were recruited via Amazon Mechanical Turk and were randomly paired (20 pairs) with one another to play as one of the agents in the Hallway (Figure 1). They were told they were playing against some other Turker. One pair was not included in the analysis due to a technical error. Subjects received \$2.00 as a base payment and a bonus of \$0.10 per round where they reached their goal (regardless of whether the other agent also reached its goal). Before being paired with another human participant, each participant went through an instruction phase with a series of “practice” grids that taught them the dynamics of the environment: arrow keys to move north, south, east or west in the grid, spacebar to wait, both agents move simultaneously, when two agents try to enter the same square, their moves fail, and the round ends once either agent reaches a goal. Example grids demonstrated outcomes in which both agents reached a goal and in which only one did and the other did not. All transitions (including transitions that didn’t involve changing location) were animated so participants could see that their action choice had registered. The instruction phase can be viewed at <http://goo.gl/SWme3n>. After the instruction phase, each human participant was paired with another human participant. Each pair played 20 rounds, which ended when either or both agents reached a goal or they had taken 30 actions without either reaching a goal.

Figure 10 plots the number of rounds an agent scored for each agent for each agent pair, which shows rich heterogeneity across participant pairs. An interface to view the actions taken by each pair of agents (and their feedback about the experiment) is available at <http://goo.gl/25IR5V>. The outcomes for a pair were encoded broadly into one of four patterns: trust (5/19), where each agent reached its goal in nearly every round; alternation (5/19), where agents reached their goals on every second round (letting the other agent reach its goal in alternating rounds); surrender (3/19), where one agent reached its goal most rounds and the other agent just got out of the way; and other (6/19), where some other pattern occurred (the majority of which were one agent reaching its goal most rounds and the other agent reaching its goal in about half of the rounds).

Contrary to the Q-learning self-play results where all pairs converged to a cooperative strategy, only about one quar-




Figure 10: Outcomes of people playing against each other in Hallway.

ter of pairs of human agents behaved cooperatively. There are two possible (and not mutually exclusive) reasons for this difference: (1) Our learning agents were given 30,000 rounds to cooperate, whereas human pairs only interacted for 20 rounds, and (2) human agents use a learning strategy that often does not result in cooperation in this environment.

To better understand whether human agents could be induced to cooperate more reliably, and to understand what happens when a learning agent is paired with a person, we ran a follow up study in which people were paired with learning algorithms.

Machine–Human Experiments

Our first study investigated pairwise interactions between humans, revealing a range of different outcomes. Our second study investigated human behavior in Hallway when pitted against reinforcement-learning agents. Our goal was to provide a baseline for studying how individuals’ behavior is influenced by the subjective preferences of a reinforcement-learning agent.

The experiment consisted of 19 participants who played against a learning agent programmed to estimate the participant’s policy by simply counting the number of times an action was taken at each state. To accelerate learning, pseudo-counts of 0.1 were added to equivalent actions in states where the participant was in the same location on the grid. The human’s estimated policy was that, at each state, the person would choose the action with the maximum count. At each turn in the experiment, the learning agent used value iteration to generate a policy against the current estimate of the human’s policy to maximize its subjective reward function. Note that the human participant was told only that they were playing another agent. The instruction phase was otherwise identical to the previous experiment.

Participants were divided into two conditions. In the selfish-opponent condition ($n = 9$), the subject played against an agent with the selfish subjective reward function.




Figure 12: Average “collisions” (turns where both agents attempt to move to the same square) plotted against total number of shared goals in the Human–Machine experiment.

In the fair-opponent condition ($n = 10$), participants played against an agent with the fair subjective reward function. As in the human–human study, each participant was paid a \$2.00 base pay and \$0.10 for each time they reached the goal. They played a total of 20 rounds that each lasted up to 30 actions each. A viewer for the results is available at <http://goo.gl/nXm6IL>.

In line with our predictions, other-regarding preferences led to changes in participant performance. In particular, fair-opponent subjects scored significantly more than selfish-opponent subjects ($t(9.0) = 2.37, p < 0.05$). Similarly, fair-opponent subjects tended to score consistently higher than selfish-opponent subjects across the experiment (Figure 11). However, the learning agents themselves scored similarly between the two conditions.

Using a similar coding scheme as in the human–human experiment, we found that games played by selfish-opponent subjects resulted in trust (5/9), surrender (1/9), and other (3/9). In contrast, games played by fair-opponent subjects resulted consistently in trust (10/10).

Broadly, the interactions between the human and Q-learning agents were either high in cooperation (many shared goals) and low in conflict (few collisions), or low in cooperation and high in conflict. The first type of interaction suggests the emergence of a “norm” that does not require either agent to explicitly defend against the other (for example by colliding). The second one corresponds to a failure to agree on a joint policy that seamlessly allows both agents to reach their goals (top-left of Figure 12). Figure 12 plots the total number of shared goals against the number of collisions averaged over all 20 rounds and shows two large clusters corresponding to the two types of interactions. However, while the fair-opponent subjects overall tend to have more shared goals, the conditions split relatively evenly between the two clusters of interactions described.

Conclusions

In this work, we showed that introducing other-regarding preferences that favored fairness to Q-learning agents re-




Figure 11: The impact of agent strategy on human-machine match ups in Hallway.

sulted in improved cooperative performance. Other multi-agent learning algorithms like Coco-Q (Sodomka et al. 2013) and Friend-Q (Littman 1994b) could also be used to promote cooperation, but will not behave as well when faced with an uncooperative opponent. In contrast, Q-learning, with subjective rewards that were different from but consistent with the selfish rewards, is able to adapt its behavior to its observed opponent.

Q-learning, even when given useful other-regarding preferences, takes many thousands of interactions to converge to good behavior. Our Amazon Mechanical Turk results showed that humans are able to converge to cooperative behavior much more quickly. We showed that a model-based approach that explicitly models its opponent’s policy could be endowed with other-regarding preferences and were able to quickly converge to good behavior on timescales comparable to people. Further, these agents were able to converge to mutually beneficial behavior when interacting with people.

Ideally, we want agents that learn, from experience, which subjective rewards are most appropriate to adopt within their population to achieve greater objective reward. Future work will investigate how to address this challenge.

One solution might be to incorporate “expert” algorithms (Crandall 2014; Megiddo and Farias 2005). Expert algorithms abstract the learning problem from learning about individual state-action pairs to learning which high-level strategy from a set of strategies to perform. In our problem, these strategies could be the set of strategies induced by different other-regarding preferences.

We reported on experiments involving (repeated) grid games played among humans, agents, and mixtures of humans and agents. As cooperation and “defection” are both aspects of our grid games, they naturally bear strong resemblances to the Prisoners’ Dilemma. In numerous experimental studies of the repeated Prisoner’s Dilemma researchers find that people cooperate more often than game theory would predict (Camerer 2003). Possible reasons for this behavior include an overall preference for cooperation, if one can determine that his/her opponent is also willing to cooperate (reliably).

As in Prisoners’ Dilemma experiments, our human exper-

iments revealed a mix of behaviors, some of which were cooperative, and others which were not. Perhaps more interesting still is the fact that there was a clear distinction in behavior among the humans who played against the agents: humans whose agent opponents were fair were able to identify and exploit this condition readily. In future, more extensive, experiments, we intend to query the participants to try to determine whether or not they were interested in and/or able to identify one-another as cooperative. If humans perceive machines as more predictable and having a preference for joint goal attainment, they may be more inclined to trust machines in joint tasks. This outcome would foster human-machine cooperation and could enable large increases in productivity by having machines share some of the workload in tasks that require more than one agent, one of whom is a human.

References

- Babes, M.; Munoz de Cote, E.; and Littman, M. L. 2008. Social reward shaping in the prisoner’s dilemma. In *7th International Joint Conference on Autonomous Agents and Multiagent Systems*, 1389–1392.
- Bowling, M., and Veloso, M. 2002. Multiagent learning using a variable learning rate. *Artificial Intelligence* 136(2):215–250.
- Bowling, M. 2000. Convergence problems of general-sum multiagent reinforcement learning. In *Proceedings of the Seventeenth International Conference on Machine Learning*, 89–94.
- Camerer, C. F. 2003. *Behavioral Game Theory: Experiments in Strategic Interaction*. Princeton University Press.
- Conitzer, V., and Sandholm, T. 2007. AWESOME: A general multiagent learning algorithm that converges in self-play and learns a best response against stationary opponents. *Machine Learning* 67:23–43.
- Crandall, J. W. 2014. Non-myopic learning in repeated stochastic games. *CoRR*.
- Gal, K.; Pfeffer, A.; Marzo, F.; and Grosz, B. J. 2004. Learning social preferences in games. In *AAAI-04*, 226–231.
- Gomes, E. R., and Kowalczyk, R. 2009. Dynamic analysis of multiagent Q-learning with e-greedy exploration. In *Pro-*

ceedings of the 2009 International Conference on Machine Learning.

Greenwald, A., and Hall, K. 2003. Correlated-Q learning. In *Proceedings of the Twentieth International Conference on Machine Learning*, 242–249.

Hu, J., and Wellman, M. P. 2003. Nash Q-learning for general-sum stochastic games. *Journal of Machine Learning Research* 4:1039–1069.

Kahneman, D.; Slovic, P.; and Tversky, A., eds. 1982. *Judgment under uncertainty: Heuristics and biases*. Cambridge University Press.

Littman, M. L., and Stone, P. 2001. Implicit negotiation in repeated games. In *Eighth International Workshop on Agent Theories, Architectures, and Languages (ATAL-2001)*, 393–404.

Littman, M. L. 1994a. Markov games as a framework for multi-agent reinforcement learning. In *Proceedings of the Eleventh International Conference on Machine Learning*, 157–163.

Littman, M. L. 1994b. Memoryless policies: Theoretical limitations and practical results. In Cliff, D.; Husbands, P.; Meyer, J.-A.; and Wilson, S. W., eds., *From Animals to Animats 3: Proceedings of the Third International Conference on Simulation of Adaptive Behavior*, 238–245. Cambridge, MA: The MIT Press.

Littman, M. L. 2001. Friend-or-foe Q-learning in general-sum games. In *Proceedings of the Eighteenth International Conference on Machine Learning*, 322–328. Morgan Kaufmann.

Megiddo, N., and Farias, D. 2005. Exploration-exploitation tradeoffs for experts algorithms in reactive environments. *Advances in Neural Information Processing Systems* 17:409–416.

Sandholm, T. W., and Crites, R. H. 1995. Multiagent reinforcement learning in the iterated prisoner's dilemma. *Biosystems* 37:144–166.

Singh, S.; Lewis, R. L.; Barto, A. G.; and Sorg, J. 2010. Intrinsically motivated reinforcement learning: An evolutionary perspective. *IEEE Transactions on Autonomous Mental Development* 2(2):70–82.

Singh, S.; Lewis, R. L.; and Barto, A. G. 2009. Where do rewards come from? In *Proceedings of the Annual Conference of the Cognitive Science Society*.

Sodomka, E.; Hilliard, E.; Littman, M.; and Greenwald, A. 2013. Coco-Q: Learning in stochastic games with side payments. *JMLR Workshop and Conference Proceedings: Proceedings of The 30th International Conference on Machine Learning* 28(3):1471–1479.

Sutton, R. S., and Barto, A. G. 1998. *Reinforcement Learning: An Introduction*. The MIT Press.

Watkins, C. J. C. H., and Dayan, P. 1992. Q-learning. *Machine Learning* 8(3):279–292.

Wunder, M.; Littman, M.; and Babes, M. 2010. Classes of multiagent Q-learning dynamics with epsilon-greedy exploration. In *Proceedings of the Twenty-Seventh International Conference on Machine Learning (ICML-10)*, 1167–1174.

Zinkevich, M.; Greenwald, A. R.; and Littman, M. L. 2005. Cyclic equilibria in Markov games. In *Advances in Neural Information Processing Systems 18*.